

CS 203: Software Tools & Techniques for AI
IIT Gandhinagar
Sem-II - 2024-25

LAB 05

Total marks: 10

Submission deadline:

Submission guidelines:

1. Code should be added to a GitHub repository, and the repository details should be shared.
2. Late submissions will be penalized 20% per day.
3. Google form submission link:

Note: By submitting this assignment solution you confirm to follow the IITGN's honor code. We shall strictly penalize the submissions containing plagiarized text/code.

Model Architecture





Figure 2. Residual learning: a building block.

Source: [Deep Residual Learning for Image Recognition, Microsoft Research](#)



Resnet 34 Network (50 is a larger version of this)

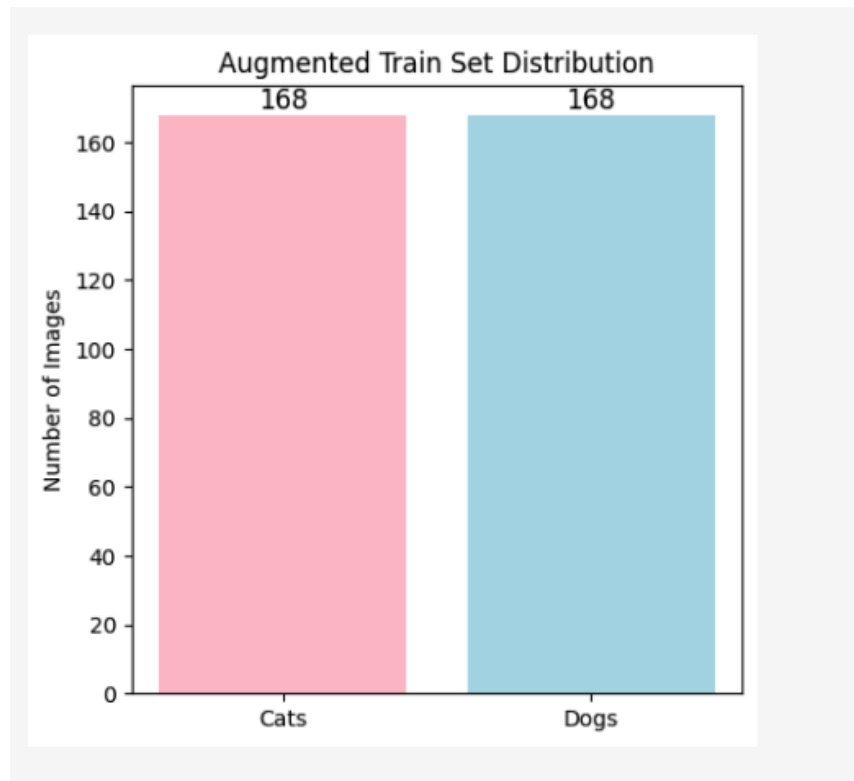
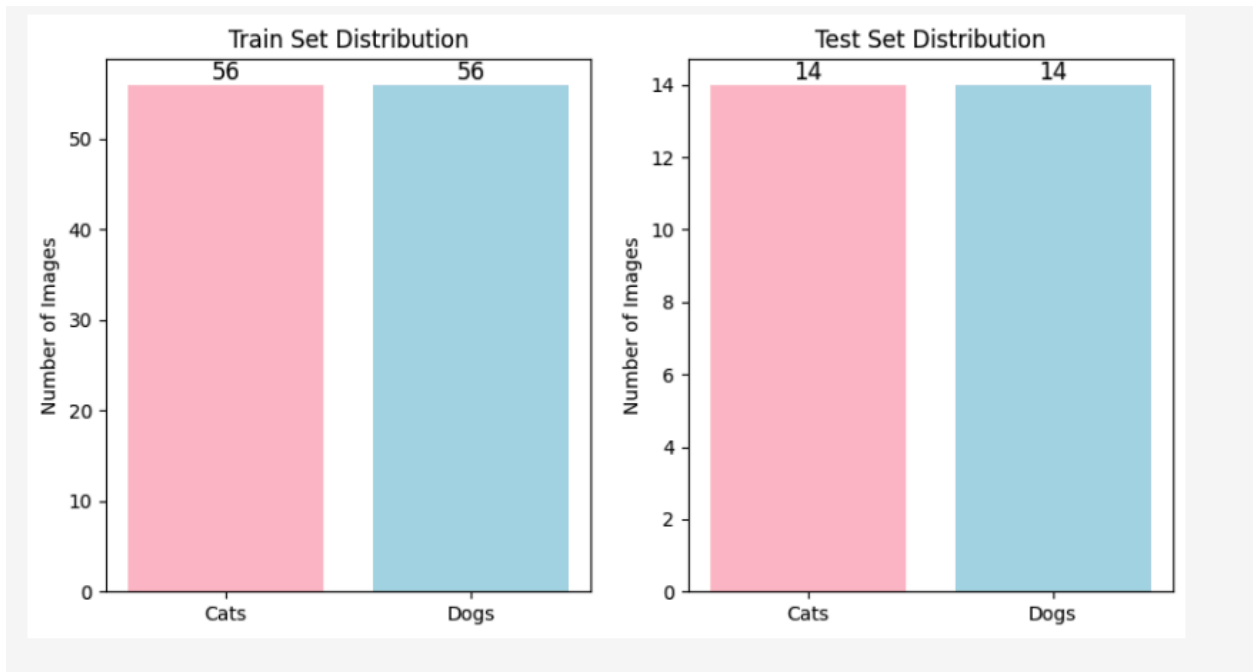
ResNet50 is a deep convolutional neural network with **50 layers**, designed for image recognition. It uses **residual connections (skip connections)** to prevent vanishing gradients and enable training of very deep networks. Pre-trained on **ImageNet**, it achieves high accuracy on various vision tasks. It's widely used in **transfer learning** for object detection, classification, and feature extraction.

Additionally, our setup uses

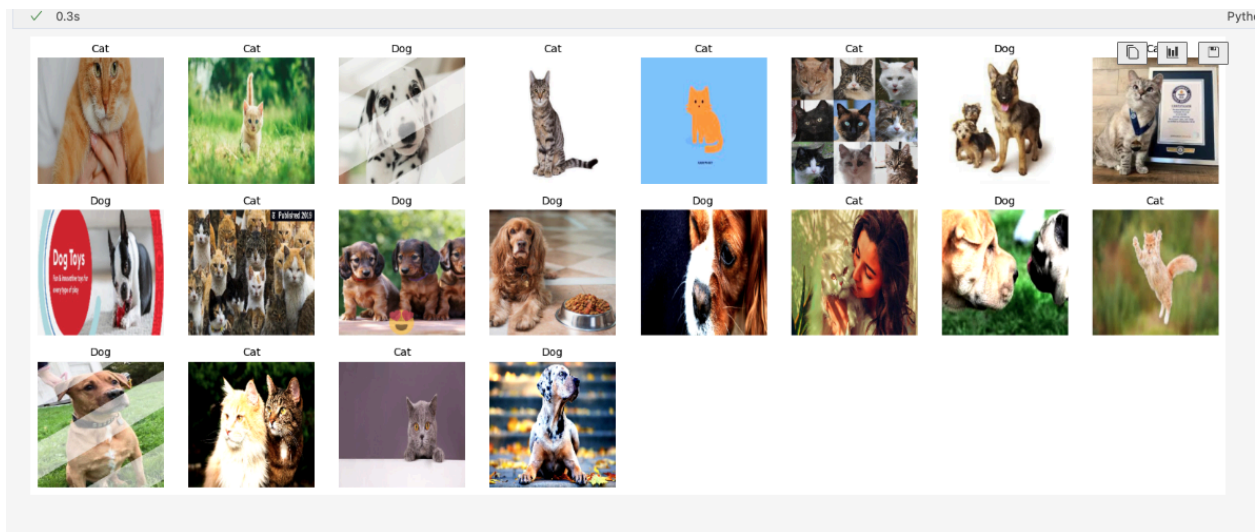
Training Parameters:

- Model: ResNet50
- Optimizer: Adam
- Learning Rate: 0.001
- Epochs: 20
- Loss Function: CrossEntropyLoss

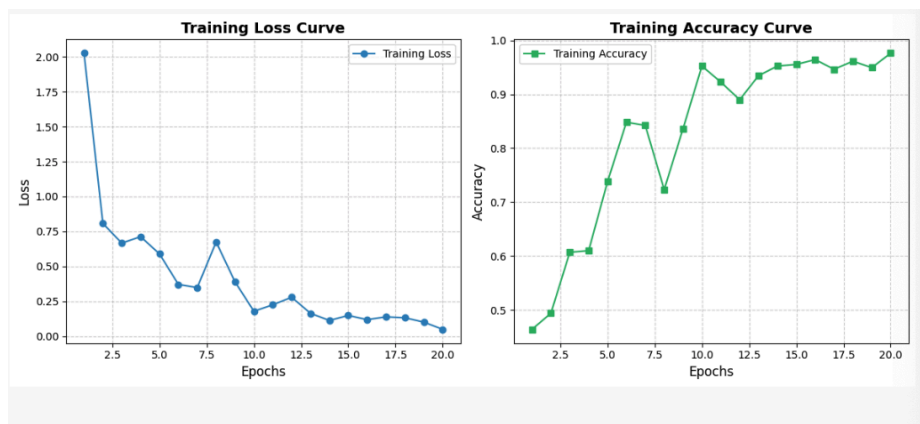
Graphs



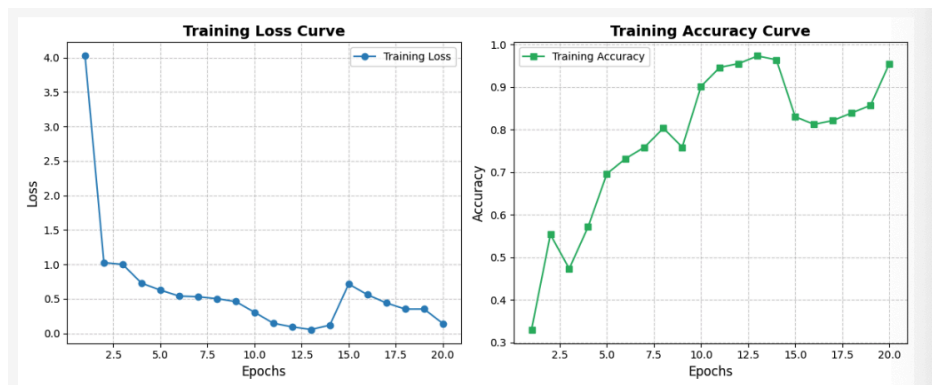
Augmented images (few)





Augmented train set training



Normal train set training



Scores

```
...  
 ResNet50 (Normal Training) Results:  
Accuracy: 0.5357  
Precision: 0.5200  
Recall: 0.9286  
F1 Score: 0.6667  
  
 ResNet50 (Augmented Training) Results:  
Accuracy: 0.6071  
Precision: 0.6000  
Recall: 0.6429  
F1 Score: 0.6207
```

Interpretation

Traintime

- We observe that the accuracy for augmented dataset rises quicker on average as compared to normal.
- This is because our augmented data offers more diversity and improves learning.

Testtime

- Here we observe, normal gives an accuracy of around 50%. This indicates that there is near to no learning that helps us, as this is binary classification and you can be right $\frac{1}{2}$ times.
- However, for the augmented dataset, our accuracy is better. And it is reasonable.
- We also feel that we could get a much higher accuracy if we used other strategies like more augmentation, validation, k-folds etc.
- For augmented dataset training, **the model generalizes better.**

Extra

- Recall drops after augmentation (from 92.86% \rightarrow 64.29%), which means the first model was aggressively predicting “dog” even when unsure. Because dog is positive class.
- F1 score is more balanced (from 66.67% \rightarrow 62.07%), suggesting the second model is more reliable overall rather than just favoring recall.
- Precision improves (from 52% \rightarrow 60%), reducing false positives.